

Binary Logistic DrPH (Epidemiology)

Kamarul Imran M

1 February 2016

Contents

Modeling Binomial Data	2
Locate files	2
Set the folder	2
Read data	2
Describe data	2
Explore data	3
Estimate parameters	3
Make inference	4
Calculate the fitted values	5
Make prediction	7
Compare models	8
Linearity in logits	10
Diagnostics for a model with a binomial response	12
Diagnostic plot	16
Interaction	16
Resources	18

Modeling Binomial Data

1. Describe data
2. Explore data - Exploratory Data Analysis
3. Estimate parameters
4. Make Inference
5. Make Prediction
6. Interpretation

Locate files

- Browse your folders.
- Look for the files.
- Check the path to the folder containing the files

Set the folder

Set our working directory. REMEMBER! your working directory (working folder) is different from my working directory

```
# this is my working directory. You have to specify yours  
setwd("E:/Epi_Stat_Matters/LectureNotes2015/binary-logistic/binary-logistic-DrPH-2015/BinaryLogisticDrPH")
```

Read data

- Read our data in the working folder
- Then, save as a csv file in our working directory

```
mydata <- read.csv("http://www.ats.ucla.edu/stat/data/binary.csv")  
write.csv2(mydata, 'logistic.csv')
```

Describe data

```
#observe data  
head(mydata, 10)
```

Rank is taken as numerical variable which does not make sense. We need to convert it to a categorical (factor) variable

```
summary(mydata)
```

##	admit	gre	gpa	rank
##	Min. :0.0000	Min. :220.0	Min. :2.260	Min. :1.000
##	1st Qu.:0.0000	1st Qu.:520.0	1st Qu.:3.130	1st Qu.:2.000
##	Median :0.0000	Median :580.0	Median :3.395	Median :2.000
##	Mean :0.3175	Mean :587.7	Mean :3.390	Mean :2.485
##	3rd Qu.:1.0000	3rd Qu.:660.0	3rd Qu.:3.670	3rd Qu.:3.000
##	Max. :1.0000	Max. :800.0	Max. :4.000	Max. :4.000

```
mydata$rank<-factor(mydata$rank)
summary(mydata$rank)
```

```
##    1    2    3    4
##  61 151 121  67
```

More fancy, we can use *psych::describe* function

```
library(psych)
describe(mydata)
```

```
##      vars   n  mean    sd median trimmed   mad   min max  range  skew
## admit     1 400  0.32  0.47   0.0   0.27  0.00   0.00  1    1.00  0.78
## gre       2 400 587.70 115.52 580.0 589.06 118.61 220.00 800 580.00 -0.14
## gpa       3 400  3.39  0.38   3.4   3.40  0.40   2.26  4    1.74 -0.21
## rank*     4 400  2.48  0.94   2.0   2.48  1.48   1.00  4    3.00  0.10
##      kurtosis  se
## admit     -1.39 0.02
## gre       -0.36 5.78
## gpa       -0.60 0.02
## rank*     -0.91 0.05
```

Explore data

Use plots like * Histogram for numerical variables * and barplot for categorical variables, at least.

Estimate parameters

- we estimate the logit or the log odds.
- We used **summary** to see the results stored as **mylogit**
- We used **coefficients** to examine the regression coefficients

```
mylogit <- glm(admit~gre+gpa+rank,family = 'binomial'(link = logit),data=mydata)
summary(mylogit)
```

```
##
## Call:
## glm(formula = admit ~ gre + gpa + rank, family = binomial(link = logit),
##      data = mydata)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6268  -0.8662  -0.6388   1.1490   2.0790
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.989979   1.139951  -3.500 0.000465 ***
## gre          0.002264   0.001094   2.070 0.038465 *
## gpa          0.804038   0.331819   2.423 0.015388 *
## rank2       -0.675443   0.316490  -2.134 0.032829 *
## rank3       -1.340204   0.345306  -3.881 0.000104 ***
## rank4       -1.551464   0.417832  -3.713 0.000205 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 458.52  on 394  degrees of freedom
## AIC: 470.52
##
## Number of Fisher Scoring iterations: 4
```

```
coefficients(mylogit)
```

```
## (Intercept)          gre          gpa          rank2          rank3
## -3.989979073  0.002264426  0.804037549 -0.675442928 -1.340203916
##          rank4
## -1.551463677
```

To obtain the odds ratios and their 95% CI, we need to exponentiate using **exp** the regression coefficients or the betas

```
exp(coefficients(mylogit))
```

```
## (Intercept)          gre          gpa          rank2          rank3          rank4
##  0.0185001  1.0022670  2.2345448  0.5089310  0.2617923  0.2119375
```

```
exp(confint(mylogit))
```

```
## Waiting for profiling to be done...
```

```
##              2.5 %      97.5 %
## (Intercept) 0.001889165 0.1665354
## gre         1.000137602 1.0044457
## gpa         1.173858216 4.3238349
## rank2       0.272289674 0.9448343
## rank3       0.131641717 0.5115181
## rank4       0.090715546 0.4706961
```

Make inference

Here, we examine the p-values (hypothesis testing) and the confidence intervals.

- First, using the method of maximum likelihood
- Next, using the SE method (function **confint.default**)

```
confint(mylogit)
```

```
## Waiting for profiling to be done...
```

```
##              2.5 %      97.5 %
## (Intercept) -6.2716202334 -1.792547080
## gre         0.0001375921  0.004435874
## gpa         0.1602959439  1.464142727
## rank2       -1.3008888002 -0.056745722
## rank3       -2.0276713127 -0.670372346
## rank4       -2.4000265384 -0.753542605
```

```
confint.default(mylogit)
```

```
##                2.5 %        97.5 %  
## (Intercept) -6.2242418514 -1.755716295  
## gre          0.0001202298  0.004408622  
## gpa          0.1536836760  1.454391423  
## rank2        -1.2957512650 -0.055134591  
## rank3        -2.0169920597 -0.663415773  
## rank4        -2.3703986294 -0.732528724
```

Calculate the fitted values

The fitted values are the expected values of the model. These expected values are the predicted probability for each observation (each patient) in the dataset

```
fitted(mylogit)
```

```
##          1          2          3          4          5          6  
## 0.17262654 0.29217496 0.73840825 0.17838461 0.11835391 0.36996994  
##          7          8          9         10         11         12  
## 0.41924616 0.21700328 0.20073518 0.51786820 0.37431440 0.40020025  
##         13         14         15         16         17         18  
## 0.72053858 0.35345462 0.69237989 0.18582508 0.33993917 0.07895335  
##         19         20         21         22         23         24  
## 0.54022772 0.57351182 0.16122101 0.43727108 0.12837525 0.19204860  
##         25         26         27         28         29         30  
## 0.43759396 0.68229503 0.57848091 0.20475422 0.42307349 0.45829857  
##         31         32         33         34         35         36  
## 0.21765393 0.28583616 0.22481919 0.42494837 0.34296523 0.21293277  
##         37         38         39         40         41         42  
## 0.48413281 0.13931720 0.26569575 0.11942769 0.18975965 0.33567002  
##         43         44         45         46         47         48  
## 0.31560404 0.17702923 0.32817441 0.18025548 0.36121718 0.11699101  
##         49         50         51         52         53         54  
## 0.07235381 0.15047417 0.31488795 0.11624726 0.23936553 0.37838478  
##         55         56         57         58         59         60  
## 0.24045684 0.39213236 0.18283980 0.10853139 0.30472142 0.12837525  
##         61         62         63         64         65         66  
## 0.33078459 0.16742893 0.28289780 0.33295972 0.30988311 0.39645173  
##         67         68         69         70         71         72  
## 0.27784995 0.51681586 0.57206626 0.69436828 0.33966212 0.07486000  
##         73         74         75         76         77         78  
## 0.15073716 0.46607599 0.24284830 0.38139149 0.20415281 0.42494837  
##         79         80         81         82         83         84  
## 0.43570986 0.65251556 0.16456653 0.31150713 0.20517359 0.08776685  
##         85         86         87         88         89         90  
## 0.21358749 0.25126279 0.34584314 0.37549461 0.55783057 0.51131037  
##         91         92         93         94         95         96  
## 0.49978497 0.63809471 0.57000341 0.26968427 0.40010880 0.37907977  
##         97         98         99        100        101        102  
## 0.22063013 0.33002244 0.31762762 0.14640896 0.11633954 0.24114689  
##        103        104        105        106        107        108  
## 0.11883427 0.28100436 0.50126183 0.35394219 0.61241920 0.25695415
```

##	109	110	111	112	113	114
##	0.11218813	0.30904921	0.17869743	0.13603549	0.10881750	0.48942091
##	115	116	117	118	119	120
##	0.35153649	0.32780508	0.29004920	0.47768876	0.68922540	0.09863460
##	121	122	123	124	125	126
##	0.38205848	0.19283124	0.13456621	0.14161529	0.35890251	0.16784107
##	127	128	129	130	131	132
##	0.55353632	0.29761787	0.29364378	0.12270194	0.32900715	0.27429792
##	133	134	135	136	137	138
##	0.35016196	0.15167362	0.26397051	0.20956391	0.16855273	0.37076538
##	139	140	141	142	143	144
##	0.37104174	0.56147017	0.48592324	0.24487554	0.27496207	0.21702497
##	145	146	147	148	149	150
##	0.18326999	0.15292361	0.30053113	0.13202601	0.36278299	0.58590453
##	151	152	153	154	155	156
##	0.69607194	0.26076336	0.48793196	0.22533437	0.27701027	0.12691355
##	157	158	159	160	161	162
##	0.20243105	0.49385024	0.40979572	0.33767745	0.31214097	0.40081797
##	163	164	165	166	167	168
##	0.44572710	0.21536268	0.33209361	0.69237989	0.12564635	0.33881603
##	169	170	171	172	173	174
##	0.27253083	0.25713529	0.16766865	0.13610230	0.27045353	0.47601029
##	175	176	177	178	179	180
##	0.17207711	0.36543032	0.20079352	0.20929210	0.22290898	0.09702710
##	181	182	183	184	185	186
##	0.29173405	0.21592659	0.53390445	0.41213948	0.10284874	0.51016205
##	187	188	189	190	191	192
##	0.23875288	0.26184001	0.28313813	0.30160149	0.29894660	0.33797096
##	193	194	195	196	197	198
##	0.29780561	0.14252603	0.37361105	0.37499458	0.20306181	0.11520619
##	199	200	201	202	203	204
##	0.25867413	0.23203530	0.29790835	0.31450637	0.69237989	0.19176895
##	205	206	207	208	209	210
##	0.62160882	0.37552455	0.62994688	0.59336886	0.17269671	0.36867073
##	211	212	213	214	215	216
##	0.23500145	0.28417171	0.21145148	0.23806753	0.39069474	0.18303592
##	217	218	219	220	221	222
##	0.29144726	0.49458858	0.36532833	0.37499458	0.18691983	0.35841190
##	223	224	225	226	227	228
##	0.38346629	0.32549498	0.37234438	0.29200523	0.40539785	0.13119209
##	229	230	231	232	233	234
##	0.30562595	0.42917277	0.17040039	0.20845157	0.25212831	0.09688336
##	235	236	237	238	239	240
##	0.65921863	0.30806878	0.40979572	0.41039144	0.10815929	0.27465027
##	241	242	243	244	245	246
##	0.19001218	0.56239934	0.19616746	0.33794240	0.41996550	0.40736827
##	247	248	249	250	251	252
##	0.39171070	0.24596016	0.29657173	0.29278619	0.20011793	0.17414395
##	253	254	255	256	257	258
##	0.43247252	0.18780755	0.26200847	0.23371984	0.30267400	0.32075797
##	259	260	261	262	263	264
##	0.33944941	0.46187255	0.34863249	0.24298996	0.16969339	0.32075797
##	265	266	267	268	269	270
##	0.26562483	0.14378335	0.15865328	0.26021896	0.41492493	0.12579904

##	271	272	273	274	275	276
##	0.48994106	0.19310678	0.45641226	0.54337733	0.27302605	0.28684953
##	277	278	279	280	281	282
##	0.22143462	0.55028996	0.16945136	0.34384116	0.49925174	0.13172559
##	283	284	285	286	287	288
##	0.21874547	0.13337693	0.28021662	0.17925207	0.60122274	0.25502619
##	289	290	291	292	293	294
##	0.23197657	0.05878643	0.38047126	0.35008696	0.46240272	0.73372225
##	295	296	297	298	299	300
##	0.29885443	0.17659931	0.45483793	0.23950580	0.34785059	0.27566478
##	301	302	303	304	305	306
##	0.36288468	0.28067279	0.22671860	0.51860565	0.07198547	0.19060160
##	307	308	309	310	311	312
##	0.44561844	0.37054412	0.28373804	0.12588934	0.30028221	0.44520022
##	313	314	315	316	317	318
##	0.30907647	0.19322270	0.17701800	0.15412239	0.18491373	0.29806393
##	319	320	321	322	323	324
##	0.18670880	0.46755914	0.14630641	0.32183935	0.12035456	0.17486941
##	325	326	327	328	329	330
##	0.12112920	0.66498227	0.38597852	0.35450549	0.33926538	0.11370930
##	331	332	333	334	335	336
##	0.39213236	0.27905234	0.34097123	0.21344965	0.20393972	0.59795326
##	337	338	339	340	341	342
##	0.16520993	0.16070084	0.45158492	0.26006097	0.14037382	0.12659514
##	343	344	345	346	347	348
##	0.22560760	0.29075910	0.18859648	0.14657301	0.35132030	0.42636137
##	349	350	351	352	353	354
##	0.25767548	0.27488628	0.57858815	0.23714608	0.18120291	0.43779599
##	355	356	357	358	359	360
##	0.40050290	0.49758253	0.38909423	0.57487559	0.25063922	0.37007654
##	361	362	363	364	365	366
##	0.59956970	0.50972425	0.35412991	0.29777892	0.49491656	0.11836196
##	367	368	369	370	371	372
##	0.12645014	0.26745319	0.63170496	0.56803162	0.39857395	0.31708679
##	373	374	375	376	377	378
##	0.37650752	0.53085361	0.41142403	0.18735742	0.41512421	0.58958954
##	379	380	381	382	383	384
##	0.20223990	0.21896113	0.46366743	0.34602886	0.34967678	0.67275941
##	385	386	387	388	389	390
##	0.18665107	0.35189341	0.52842881	0.34287938	0.33908140	0.40275050
##	391	392	393	394	395	396
##	0.40093595	0.48719398	0.22202911	0.43872524	0.25342327	0.48866999
##	397	398	399	400		
##	0.16550430	0.18106222	0.46366743	0.30073055		

Make prediction

Similarly, one of the important objectives in modelling is to perform prediction based on the model using new data.

We can perform these predictions:

1. Predict the log odds for having the outcome

2. Predict the conditional probability for having the outcome

Let us say we have these data

gre = 380 gpa = 3.61 rank = 3

First, we create a data frame

```
new_data1 <- data.frame( gre = 380, gpa = 3.61, rank = '3')
```

Now, we can perform the prediction

```
pred.logit<-predict(mylogit,newdata = new_data1, type='link')
pred.logit
```

```
##          1
## -1.567126
```

We can confirm this by calculate this

```
-3.99+0.00226*380+0.8041*3.61-1.34
```

```
## [1] -1.568399
```

Notice, that similarity between **predict(x, type='response')** and **fitted** Remember, we can calculate the conditional probability of having the outcome

```
pred.prob<-predict(mylogit, newdata = new_data1, type='response')
pred.prob
```

```
##          1
## 0.1726265
```

We can verify this by two ways

```
head(fitted(mylogit))
```

```
##          1          2          3          4          5          6
## 0.1726265 0.2921750 0.7384082 0.1783846 0.1183539 0.3699699
```

```
# calculate the logistic probability
exp(-1.567)/(1+exp(-1.567))
```

```
## [1] 0.1726445
```

Compare models

We compare a model with **vs** and without **gre**. This is done using the deviance

```
summary(mylogit)
```

```
##
## Call:
## glm(formula = admit ~ gre + gpa + rank, family = binomial(link = logit),
##      data = mydata)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6268  -0.8662  -0.6388   1.1490   2.0790
##
## Coefficients:
```



```

##               Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.989979   1.139951  -3.500 0.000465 ***
## gre          0.002264   0.001094   2.070 0.038465 *
## gpa          0.804038   0.331819   2.423 0.015388 *
## rank2        -0.675443   0.316490  -2.134 0.032829 *
## rank3        -1.340204   0.345306  -3.881 0.000104 ***
## rank4        -1.551464   0.417832  -3.713 0.000205 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 458.52  on 394  degrees of freedom
## AIC: 470.52
##
## Number of Fisher Scoring iterations: 4
mylogit2 <- glm(admit~gpa+rank,family = 'binomial'(link = logit),data=mydata)
summary(mylogit2)

##
## Call:
## glm(formula = admit ~ gpa + rank, family = binomial(link = logit),
##      data = mydata)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5055  -0.8663  -0.6590   1.1505   2.0913
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -3.4636     1.1003  -3.148 0.001645 **
## gpa           1.0521     0.3102   3.392 0.000694 ***
## rank2        -0.6810     0.3141  -2.168 0.030181 *
## rank3        -1.3919     0.3419  -4.071 4.68e-05 ***
## rank4        -1.5943     0.4152  -3.840 0.000123 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 462.88  on 395  degrees of freedom
## AIC: 472.88
##
## Number of Fisher Scoring iterations: 4
anova(mylogit,mylogit2,test = 'Chisq')

## Analysis of Deviance Table
##
## Model 1: admit ~ gre + gpa + rank
## Model 2: admit ~ gpa + rank
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)

```

```
## 1      394      458.52
## 2      395      462.88 -1  -4.3578  0.03684 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Linearity in logits

gre is tested for linearity in logit. **gre** is linear but it is rescaled to produce less decimals

The linearity of logits is tested using library **mfp**

```
library(mfp)
```

```
## Loading required package: survival
```

```
mylogit3 <- mfp(admit~fp(gre)+gpa+rank,family = 'binomial'(link = logit),data=mydata,verbose=T)
```

```
##
## Variable      Deviance      Power(s)
## -----
## Cycle 1
## rank2
##           463.096
##           458.517      1
##
##
## rank3
##           474.043
##           458.517      1
##
##
## rank4
##           473.551
##           458.517      1
##
##
## gpa
##           464.532
##           458.517      1
##
##
## gre
##           462.875
##           458.517      1
##           458.415      0
##           458.366     -2 -2
##
##
## Transformation
##           shift scale
```

```

## rank2      0      1
## rank3      0      1
## rank4      0      1
## gpa        0      1
## gre        0 1000
##
## Fractional polynomials
##      df.initial select alpha df.final power1 power2
## rank2      1      1 0.05      1      1      .
## rank3      1      1 0.05      1      1      .
## rank4      1      1 0.05      1      1      .
## gpa        1      1 0.05      1      1      .
## gre        4      1 0.05      1      1      .
##
##
## Transformations of covariates:
##      formula
## gre  I((gre/1000)^1)
## gpa      gpa
## rank      rank
##
##
## Deviance table:
##      Resid. Dev
## Null model    499.9765
## Linear model  458.5175
## Final model   458.5175
summary(mylogit3)

##
## Call:
## glm(formula = admit ~ rank + gpa + I((gre/1000)^1), family = binomial(link = logit),
##      data = mydata)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6268  -0.8662  -0.6388   1.1490   2.0790
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -3.9900     1.1400  -3.500 0.000465 ***
## rank2         -0.6754     0.3165  -2.134 0.032829 *
## rank3         -1.3402     0.3453  -3.881 0.000104 ***
## rank4         -1.5515     0.4178  -3.713 0.000205 ***
## gpa           0.8040     0.3318   2.423 0.015388 *
## I((gre/1000)^1) 2.2644     1.0940   2.070 0.038465 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 458.52  on 394  degrees of freedom
## AIC: 470.52

```

```
##
## Number of Fisher Scoring iterations: 4
mylogit3$fptable

##      df.initial select alpha df.final power1 power2
## rank2          1      1 0.05          1      1      .
## rank3          1      1 0.05          1      1      .
## rank4          1      1 0.05          1      1      .
## gpa            1      1 0.05          1      1      .
## gre            4      1 0.05          1      1      .
```

Diagnostics for a model with a binomial response

To do these diagnostics, you need to load `library('LogisticDx')`.

First, we produce the diagnostic measures for a binary regression model by covariate pattern

Next, we produce the Goodness-of-fit for binomial regression. Usually, the number of groups (quantiles) equal 10 to perform the Hosmer-Lemeshow test. At the same time, we plot the ROC curve


Similarly, we can check the **auc** value

```
library('LogisticDx')
dx(mylogit2,byCov=T)
```

```
##      (Intercept)  gpa rank2 rank3 rank4 y      P n      yhat
## 1:          1 3.94      1      0      0 1 0.5001242 2 1.0002483
## 2:          1 3.30      1      0      0 1 0.3378663 3 1.0135990
## 3:          1 3.99      0      1      0 1 0.3412259 3 1.0236777
## 4:          1 3.35      1      0      0 1 0.3497327 3 1.0491980
## 5:          1 3.17      1      0      0 1 0.3079792 3 0.9239375
## ---
## 256:          1 3.52      0      0      1 2 0.2051122 3 0.6153366
## 257:          1 2.68      0      1      0 1 0.1154721 1 0.1154721
## 258:          1 3.00      0      0      1 1 0.1299149 1 0.1299149
## 259:          1 2.42      0      0      0 1 0.2854368 1 0.2854368
## 260:          1 2.65      0      1      0 1 0.1122874 1 0.1122874
##      Pr      dr      h      sPr      sdr
## 1: -0.0003511985 -0.0003511985 0.014377039 -0.0003537507 -0.0003537506
## 2: -0.0165997514 -0.0166181074 0.006902268 -0.0166573376 -0.0166757573
## 3: -0.0288329487 -0.0288874508 0.015618469 -0.0290607843 -0.0291157171
## 4: -0.0595625368 -0.0597856465 0.006844794 -0.0597674360 -0.0599913131
## 5: 0.0951239167 0.0944359041 0.007474218 0.0954814103 0.0947908121
## ---
## 256: 1.9798618855 1.7253451406 0.017098521 1.9970084425 1.7402874601
## 257: 2.7676886108 2.0778480687 0.011720502 2.7840519420 2.0901328741
## 258: 2.5879228606 2.0203343433 0.013448875 2.6055027084 2.0340585431
## 259: 1.5822145401 1.5834991249 0.034641074 1.6103525417 1.6116599714
## 260: 2.8117113691 2.0912646219 0.011927131 2.8286306943 2.1038486967
##      dChisq      dDev      dBhat
## 1: 1.251395e-07 1.251395e-07 1.825379e-09
## 2: 2.774669e-04 2.780809e-04 1.928462e-06
## 3: 8.445292e-04 8.477250e-04 1.339953e-05
## 4: 3.572146e-03 3.598958e-03 2.461912e-05
```

```
## 5: 9.116700e-03 8.985298e-03 6.865333e-05
## ---
## 256: 3.988043e+00 3.028600e+00 6.937586e-02
## 257: 7.750945e+00 4.368655e+00 9.192235e-02
## 258: 6.788644e+00 4.137394e+00 9.254425e-02
## 259: 2.593235e+00 2.597448e+00 9.305602e-02
## 260: 8.001152e+00 4.426179e+00 9.658274e-02
```

```
fit.mylogit2<-gof(mylogit2,g=10,plotROC = T)
```



```
fit.mylogit2
```

```
##      chiSq df    pVal
## PrI 400.78 395 0.409645
## drI 462.88 395 0.010379 *
## PrG 255.36 255 0.481870
## drG 294.60 255 0.044609 *
## PrCT 255.36 255 0.481870
## drCT 294.60 255 0.044609 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##      val df    pVal
## HL chiSq      8.190458 8 0.415090
## mHL F         1.484491 9 0.153918
## OsRo Z        0.027903 NA 0.977740
## SstPgeq0.5 Z   0.217564 NA 0.827769
## SstP10.5 Z     0.422859 NA 0.672398
```

```
## SstBoth chiSq    0.226144  2 0.893086
## SllPgeq0.5 chiSq 0.047739  1 0.827045
## SllPl0.5 chiSq   0.176745  1 0.674186
## SllBoth chiSq    0.177398  2 0.915121

#area under curve
fit.mylogit2$auc

##          auc lower 95% CI upper 95% CI
## 68.01073    62.39737    73.62409
## attr("interpret")
## [1] "auc = 0.5          --> useless"    "0.7 < auc < 0.8 --> good"
## [3] "0.8 < auc < 0.9 --> excellent"

#chi square test for gof
fit.mylogit2$chiSq

##      test    chiSq df      pVal
## 1: PrI 400.7833 395 0.40964543
## 2: drI 462.8753 395 0.01037867
## 3: PrG 255.3601 255 0.48187033
## 4: drG 294.6050 255 0.04460896
## 5: PrCT 255.3601 255 0.48187033
## 6: drCT 294.6050 255 0.04460896

#contingency table for HL test
fit.mylogit2$ctHL

##      P y1      y1hat y0      y0hat  n      Pbar
## 1: 0.149  6  5.037833 34 34.96217 40 0.1259458
## 2: 0.192  5  6.824095 35 33.17591 40 0.1706024
## 3: 0.219  7  8.260683 33 31.73932 40 0.2065171
## 4: 0.258 11  9.327547 28 29.67245 39 0.2391679
## 5: 0.293 12 11.338710 29 29.66129 41 0.2765539
## 6: 0.34  15 12.556219 25 27.44378 40 0.3139055
## 7: 0.376 14 13.796927 25 25.20307 39 0.3537674
## 8: 0.437 14 16.561867 27 24.43813 41 0.4039480
## 9: 0.511 13 18.253216 25 19.74678 38 0.4803478
## 10: 0.678 30 25.042903 12 16.95710 42 0.5962596

#GOF test
fit.mylogit2$gof


##      test  stat      val df      pVal
## 1:      HL chiSq 8.19045835  8 0.4150903
## 2:      mHL      F 1.48449073  9 0.1539182
## 3:      OsRo      Z 0.02790263 NA 0.9777398
## 4: SstPgeq0.5      Z 0.21756392 NA 0.8277689
## 5: SstPl0.5      Z 0.42285933 NA 0.6723979
## 6: SstBoth chiSq 0.22614408  2 0.8930863
## 7: SllPgeq0.5 chiSq 0.04773921  1 0.8270450
## 8: SllPl0.5 chiSq 0.17674490  1 0.6741857
## 9: SllBoth chiSq 0.17739848  2 0.9151208
```

We can also perform model fitness by using *ROCR* package

```
library(ROCR)
```

```
## Loading required package: gplots
##
## Attaching package: 'gplots'
## The following object is masked from 'package:stats':
##
##      lowess
pred.prob2<-predict(mylogit2, type='response')
head(pred.prob2)

##          1          2          3          4          5          6
## 0.2577653 0.2700256 0.6779993 0.1542276 0.1218149 0.2712217
pred.prob22<-prediction(pred.prob2, mydata$admit )
pred.prob22f<-performance(pred.prob22, measure='tpr', x.measure='fpr')
plot(pred.prob22f)
```



Using *ROCR* package, we can also calculate the *AUC*

```
auc2<-performance(pred.prob22, measure='auc')
auc2@y.values[[1]]
```

```
## [1] 0.6801073
```

Another package is *MKmisc* package, to perform the Hosmer-Lemeshow test of GOF

```
library(MKmisc)
```

```
##
```

```
## Attaching package: 'MKmisc'
## The following object is masked from 'package:psych':
##
##      corPlot
HLgof.test(fit = fitted(mylogit2), obs = mydata$admit)

## $C
##
## Hosmer-Lemeshow C statistic
##
## data: fitted(mylogit2) and mydata$admit
## X-squared = 9.8564, df = 8, p-value = 0.2752
##
##
## $H
##
## Hosmer-Lemeshow H statistic
##
## data: fitted(mylogit2) and mydata$admit
## X-squared = 7.6802, df = 8, p-value = 0.4653
```

Diagnostic plot

Will return many diagnostic plot

```
plot(mylogit2)
```

Using K-fold validation. Will not discuss here.

Interaction

Let use see how we deal an interaction. First, read data from this text file.

Columns (variables) no 2, and from 5 to 10 need to be converted to categorical (factor) variables

```
data.1<-read.table("LOWBWT.txt",header=T)
data.1[,c(2,5:10)]<-lapply(data.1[,c(2,5)],factor)
```

To simulate a binary predictor variable, we now recode LWT to LWD (LWT<110 vs >=110)

```
data.1$LWD<-findInterval(data.1$LWT,110)
data.1$LWD<-factor(data.1$LWD,labels = c("less 110", ">=110"))
head(data.1$LWD,10)
```

```
## [1] >=110    >=110    >=110    less 110 less 110 >=110    less 110
## [8] >=110    >=110    >=110
## Levels: less 110 >=110
```

```
head(data.1$LWT,10)
```

```
## [1] 120 130 187 105 85 150 97 128 132 165
```

```
str(data.1$LWD)
```



```
## Factor w/ 2 levels "less 110", ">=110": 2 2 2 1 1 2 1 2 2 2 ...
```

Model the relationship; outcome (LOW=0,1) with predictors of LWD and AGE interact with each other. You may try with using #

```
mod.lwd.age<-glm(LOW~LWD*AGE,family = binomial(link =logit ),data=data.1)
summary(mod.lwd.age)
```

```
##
## Call:
## glm(formula = LOW ~ LWD * AGE, family = binomial(link = logit),
##      data = data.1)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4257  -0.8554  -0.6960   1.1602   2.0329
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.16959    1.46515  -0.798   0.4247
## LWD>=110       1.94409    1.72481   1.127   0.2597
## AGE           0.05262    0.06449   0.816   0.4145
## LWD>=110:AGE  -0.13220    0.07570  -1.746   0.0807 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 234.67  on 188  degrees of freedom
## Residual deviance: 221.14  on 185  degrees of freedom
## AIC: 229.14
##
## Number of Fisher Scoring iterations: 4
```

Predict our model using new data. Before doing so, we need to create a dataset containing new data

```
newdata.2<-data.frame(AGE=c(15,15,20,20),LWD=rep(c("less 110", ">=110"),2))
```

Now let us predict the log odds

```
predict(mod.lwd.age,newdata=newdata.2)
```

```
##           1           2           3           4
## -0.3802252 -0.4190874 -0.1171023 -0.8169483
```

```
newdata.2
```

```
##   AGE    LWD
## 1  15 less 110
## 2  15  >=110
## 3  20 less 110
## 4  20  >=110
```

Can you prove these?

```
-1.1696+0.0526*15+0
```

```
## [1] -0.3806
```

```
-1.1696+1.944*1+0.0526*15-0.1322*1*15
```

```
## [1] -0.4196
```

```
-1.1696+0+0.0526*20+0
```

```
## [1] -0.1176
```

```
-1.1696+1.944*1+0.0526*20-0.1322*1*20
```

```
## [1] -0.8176
```

Resources

1. <http://www.ats.ucla.edu/stat/r/dae/logit.htm>
2. https://cran.r-project.org/web/packages/HSAUR/vignettes/Ch_logistic_regression_glm.pdf